

## What deep learning is good for



- **Problems with long lists of rules**—when the traditional approach fails, machine learning/deep learning may help.
- **Continually changing environments**—deep learning can adapt ('learn') to new scenarios.
- **Discovering insights within large collections of data**—can you imagine trying to hand-craft rules for what 101 different kinds of food look like?

## What deep learning is <sup>(typically)</sup> not good for



- **When you need explainability**—the patterns learned by a deep learning model are typically uninterpretable by a human.
- **When the traditional approach is a better option** — if you can accomplish what you need with a simple rule-based system.
- **When errors are unacceptable** — since the outputs of deep learning model aren't always predictable.
- **When you don't have much data** — deep learning models usually require a fairly large amount of data to produce great results.

# Machine Learning vs. Deep Learning

(common algorithms)

- Random forest
- Gradient boosted models
- Naive Bayes
- Nearest neighbour
- Support vector machine
- ...many more

(since the advent of deep learning these are often referred to as "shallow algorithms")

- Neural networks
- Fully connected neural network
- Convolutional neural network
- Recurrent neural network
- Transformer
- ...many more

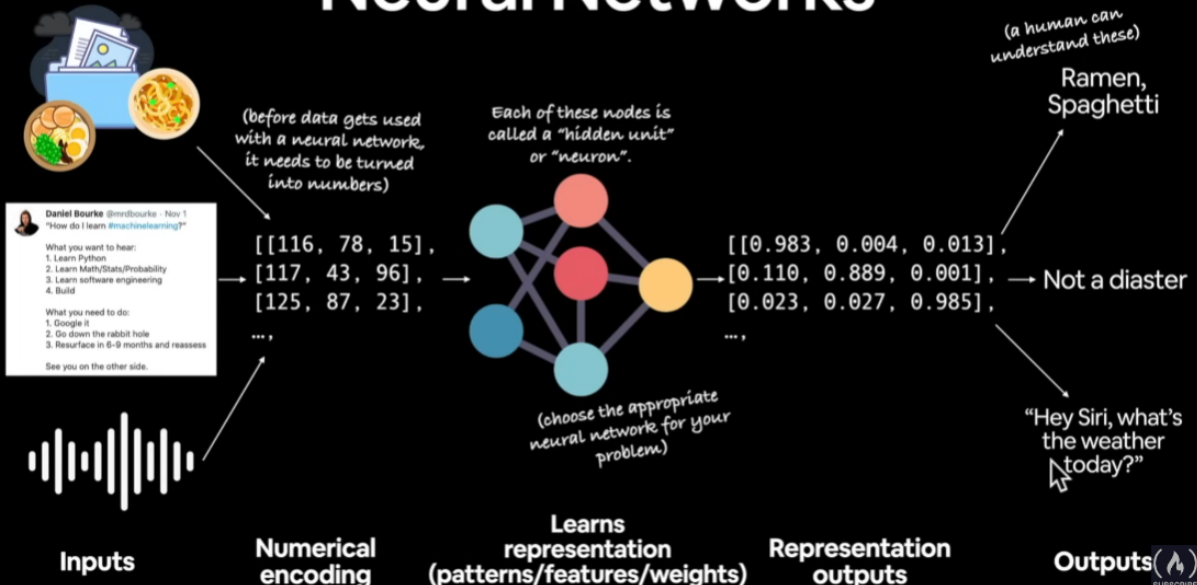
What we're focused on building (with PyTorch)

(depending how you represent your problem, many algorithms can be used for both)

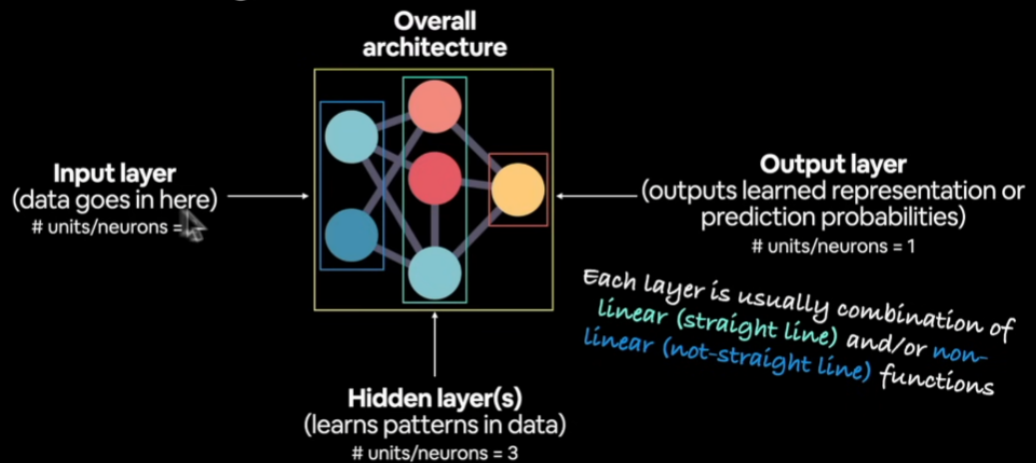
Structured data

Unstructured data

## Neural Networks



# Anatomy of Neural Networks



**Note:** “patterns” is an arbitrary term, you’ll often hear “embedding”, “weights”, “feature representation”, “feature vectors” all referring to similar things.

## Types of Learning



**Supervised Learning**



**Unsupervised & Self-supervised Learning**



**Transfer Learning**

We’ll be writing code to do these,  
but the style of code can be adopted across learning paradigms.

# What we're going to cover

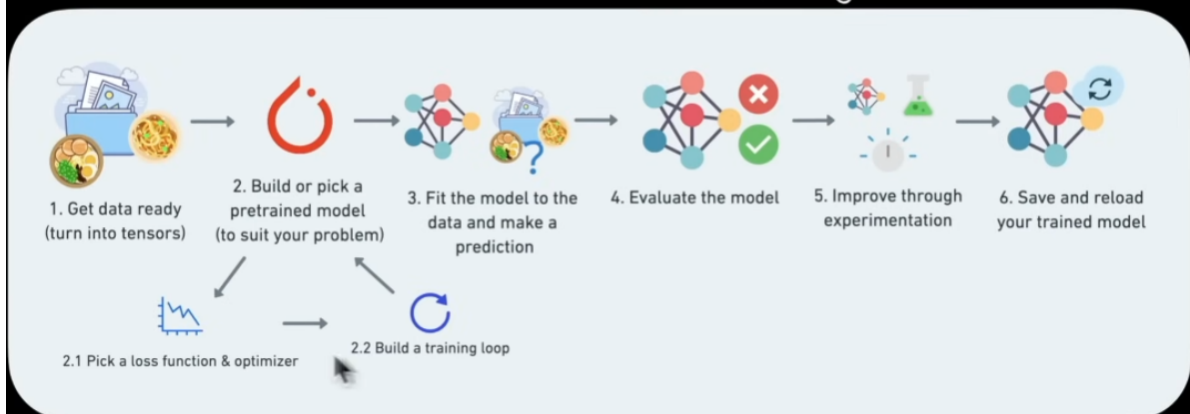
(broadly)

- Now:
  - PyTorch basics & fundamentals (dealing with tensors and tensor operations)
- Later:
  - Preprocessing data (getting it into tensors)
  - Building and using pretrained deep learning models
  - Fitting a model to the data (learning patterns)
  - Making predictions with a model (using patterns)
  - Evaluating model predictions
  - Saving and loading models
  - Using a trained model to make predictions on custom data

# What we're going to cover

## A PyTorch workflow

(one of many)



# How to approach this course

```
1 # 1. Construct a model class that subclasses nn.Module
2 class CircleModelV0(nn.Module):
3     def __init__(self):
4         super().__init__()
5         # 2. Create 2 nn.Linear layers
6         self.layer_1 = nn.Linear(in_features=2, out_features=5)
7         self.layer_2 = nn.Linear(in_features=5, out_features=1)
8
9     # 3. Define a forward method containing the forward pass computation
10    def forward(self, x):
11        # Pass the data through both layers
12        return self.layer_2(self.layer_1(x))
13
14 # 4. Create an instance of the model and send it to target device
15 model_0 = CircleModelV0().to(device)
16 model_0
```

## 1. Code along

Motto #1: *if in doubt, run the code!*



(including the  
"dumb" ones)

## 4. Ask questions

Motto #2:  
*Experiment, experiment,  
experiment!*



## 2. Explore and experiment

Motto #3:  
*visualize, visualize, visualize!*



## 3. Visualize what you don't understand



## 5. Do the exercises



## 6. Share your work

